

Artificial Intelligence

Lec 1 - Introduction

Pratik Mazumder

Timings

Classroom: LHB 105

Lectures:

Tue: 10am-11am

Wed: 9am-10am

Fri: 9am-10am

Office Timings:

CSE 223, Thursday 11am-12pm

Syllabus

Introduction: Uninformed search strategies, Greedy best-first search, And-Or search, Uniform cost search, A* search, Memory-bounded heuristic search, Local and evolutionary searches (9 Lectures)
Constraint Satisfaction Problems: Backtracking search for CSPs, Local search for CSPs (3 Lectures)
Adversarial Search: Optimal Decision in Games, The minimax algorithm, Alpha-Beta pruning, Expectimax search (4 Lectures)

Knowledge and Reasoning: Propositional Logic, Reasoning Patterns in propositional logic; First order logic: syntax, semantics, Inference in First order logic, unification and lifting, backward chaining, resolution (9 Lectures)

Planning: Situation Calculus, Deductive planning, STRIPES, sub-goal, Partial order planner (3 Lectures)

Bayesian Network, Causality, and Uncertain Reasoning: Probabilistic models, directed and undirected models, inferencing, causality, Introduction to Probabilistic reasoning (6 lectures)

Reinforcement Learning: MDP, Policy, Q-value, Passive RL, Active RL, Policy Search (8 Lectures)

Textbook: Russel, S., and Norvig, P., (2015), Artificial Intelligence: A Modern Approach, 3rd Edition, Prentice Hall

Course Logistics

Assignments: 2 [30%]

Quizzes: 2 [10%]

Minor: 20%

Major: 40%

Optional: Project on Reinforcement Learning - Certificates for Top 3 Teams

Textbook: Russel, S., and Norvig, P., (2015), Artificial Intelligence: A Modern Approach, 3rd Edition, Prentice Hall

Google Classroom for Assignment Submission, Announcements: Invite Sent to Registered Students

Lecture Site: <https://sites.google.com/iitj.ac.in/ai2024>

Doubts?


Intelligence

- It is a complex and multifaceted concept that generally refers to the ability to learn, understand, reason, solve problems, and adapt to new situations.
- Human intelligence is the intellectual capability of humans.
- Howard Gardner's Multiple Intelligences
 - In order to capture the full range of abilities and talents that people possess, Gardner theorizes that people do not have just an intellectual capacity,
 - but have many kinds of intelligence, including musical, interpersonal, spatial-visual, and linguistic intelligences
- In this course, we deal with intelligence required to solve problems

Artificial Intelligence


- Human beings have tried to **understand** how we think, perceive, understand, predict, and manipulate the complicated world around us.
- The field of Artificial Intelligence (AI) deals with **building intelligent entities/agents** which can perceive, understand, predict, and manipulate the world like us and may be even better than us.
- Performing a Task vs Intelligently Performing a Task

Let's Play a Turn-based Game




Let's Play a Turn-based Game

Task: Player should get 3 of his/her marks in marks in a row/column/diagonal




Let's Play a Turn-based Game

Task: Player should get 3 of his/her marks in marks in a row/column/diagonal




Let's Play a Turn-based Game

Task: Player should get 3 of his/her marks in marks in a row/column/diagonal




Let's Play a Turn-based Game

Task: Player should get 3 of his/her marks in marks in a row/column/diagonal




Let's Play a Turn-based Game

Task: Player should get 3 of his/her marks in marks in a row/column/diagonal



Let's Play a Turn-based Game

Task: Player should get 3 of his/her marks in marks in a row/column/diagonal




Why not this step instead?



Let's Play a Turn-based Game

Task: Player should get 3 of his/her marks in marks in a row/column/diagonal




Why not this step instead?



Let's Play a Turn-based Game

Task: Player should get 3 of his/her marks in marks in a row/column/diagonal




Why not this step instead?



Let's Play a Turn-based Game

Goal: To be the **First** Player to get 3 of his/her marks in marks in a row/column/diagonal



Why not this step instead?


Simply completing a task is not enough.

Intelligently solving a task involves considering the rewards/penalties



Let's Play a Turn-based Game

Goal: To be the **First** Player to get 3 of his/her marks in marks in a row/column/diagonal



Why not this step instead?

Simply completing a task is not enough.

Intelligently solving a task involves considering the rewards/penalties

Isn't that how we learn too? – Parents, Teachers

Artificial Intelligence

- How do you consider anything intelligent: **Turing Test (originally called the imitation game)**
 - was designed to provide a satisfactory operational definition of intelligence
 - A computer passes the test if a human interrogator, after posing some written questions, cannot tell whether the written responses come from a person or from a computer.

Artificial Intelligence



- In order to pass the Turing Test, the entity should possess the following capabilities:
 - **natural language processing** to enable it to communicate successfully
 - **knowledge representation** to store what it knows or hears
 - **automated reasoning** to use the stored information to answer questions and to draw new conclusion
 - **machine learning** to adapt to new circumstances and to detect and extrapolate patterns.
- Not a fool-proof measure

Artificial Intelligence

- Turing's test deliberately avoided direct physical interaction between the interrogator and the entity or computer.
- The Total Turing Test includes a video signal so that the interrogator can test the subject's perceptual abilities, as well as the opportunity for the interrogator to pass physical objects "through the hatch".
- To pass the total Turing Test, the computer will need
 - **computer vision** to perceive objects
 - **robotics** to manipulate objects and move about.



Artificial Intelligence: Agent

- An agent is an entity that perceives and acts.
- Operate **autonomously**, perceive their environment, persist over a prolonged time period, adapt to change, and create and pursue goals.
- Building AI for a task = Agent





Designing Rational Agent

- A **rational** agent is one that acts so as to achieve the best outcome or, when there is uncertainty, the best expected outcome against an ideal performance measure
- System is rational if it does the “right thing”, given what it knows, e.g., decisions taken by an autonomous car
- A rational agent selects actions that maximize its (expected) utility.
- Characteristics of the **percepts**, **environment**, and **action space** dictate techniques for selecting rational actions





Rational Decisions

- Rational: Achieve goals while Maximizing pre-defined rewards/utilities
- Rationality only concerns what decisions are made (not the thought process behind them)
- Being rational means maximizing your expected utility






Reflex Agent

- Choose action based on current percept (and maybe memory)
- May have memory or a model of the world's current state
- Do not consider the future consequences of their actions
- Not planning ahead



Reflex Agent


- Choose action based on current percept (and maybe memory)
- May have memory or a model of the world's current state
- Do not consider the future consequences of their actions
- Not planning ahead



Stuck


Planning Agent

- Planning agents:
 - Ask “what if”
 - Decisions based on (hypothesized) consequences of actions
 - Must have a model of how the world evolves in response to actions
 - Must formulate a goal (test)
- Optimal vs. complete planning:
 - Optimal: Find best solution
 - Complete: If there is a solution then find it, even if it is not the best




Search Problems

- Framework for solving problems
- A search problem consists of:
 - A state space: List of all possible states, i.e., all possible configurations of elements/conditions in the world or environment, e.g., passing one pen across the class, situations in a pac man game.
 - A successor function (with actions, costs): Action may also change the state.
 - A start state
 - A goal test: There can be multiple goal states
- A solution is a sequence of actions (a plan) which transforms the start state to a goal state



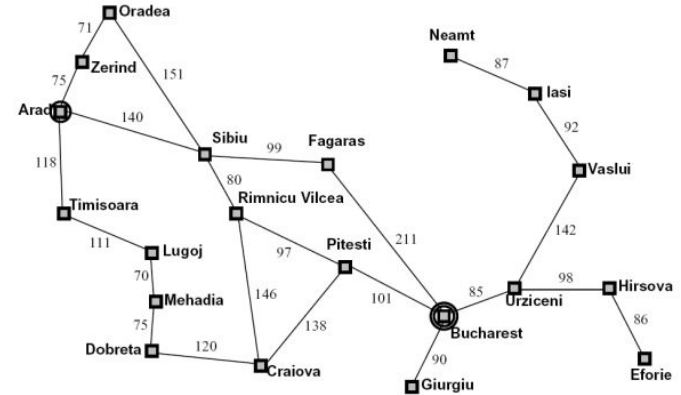
Search Problems

- There are algorithms to solve search problems
- Learn to map real-world problems to search problems
 - Then the above algorithms can give you a way to solve the problem




Search Problems: Example

- State space:
 - Cities
- Successor function:
 - Roads: Go to adjacent city with cost = distance
- Start state:
 - Arad
- Goal test:
 - Is state == Bucharest?
- Solution?



State Space

The **world state** includes every last detail of the environment



A **search state** keeps only the details needed for planning (abstraction)

■ Problem: Pathing


- States: (x,y) location
- Actions: NSEW
- Successor: update location only
- Goal test: is $(x,y)=END$

■ Problem: Eat-All-Dots

- States: $\{(x,y), \text{dot booleans}\}$
- Actions: NSEW
- Successor: update location and possibly a dot boolean
- Goal test: dots all false



State Space

- World state:
 - Agent positions: 120
 - Food count: 30
 - Ghost positions: 12
 - Agent facing: NSEW
- How many
 - World states?
 - $120 \times (2^{30}) \times (12^2) \times 4$
- States for pathing?
 - 120
- States for eat-all-dots?
 - $120 \times (2^{30})$




State Space Graphs

- A mathematical representation of a search problem
 - **Nodes** are (abstracted) world configurations/states
 - **Arcs** represent successors (action results)
 - The **goal test** is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!




Search Trees

- A search tree:
 - A “what if” tree of plans and their outcomes
 - The start state is the root node
 - Children correspond to successors
 - Nodes show states, but correspond to plans that achieve those states



State Space Graphs vs. Search Trees


State Space Graph



*Each NODE in
the search tree is
an entire PATH in
the state space
graph.*

*We construct both
on demand – and
we construct as
little as possible.*

Search Tree



Try: Represent Problems/Tasks as Search Problems

- Cook a Meal with some Ingredients
- Build a Team Project
- Build an agent to play Call of Duty