


Artificial Intelligence

Lec 7: Local Search (contd.)

Pratik Mazumder

Local Search: Hill-climbing (Greedy Local Search)

- Suppose you are at a state with multiple neighbors.
- For each neighbor, you have the objective function value.
- Strategy:
 - You can pick the neighbor with the best objective function and repeat
 - But then, when will you stop?
 - If **none** of the neighbours have a **lower objective function value** then stop [Case: when objective function has to be minimized]
 - If **none** of the neighbours have a **higher objective function value** then stop [Case: when objective function has to be maximized]
- Hill-climbing or Greedy Local Search



Hill-climbing (Greedy Local Search)

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

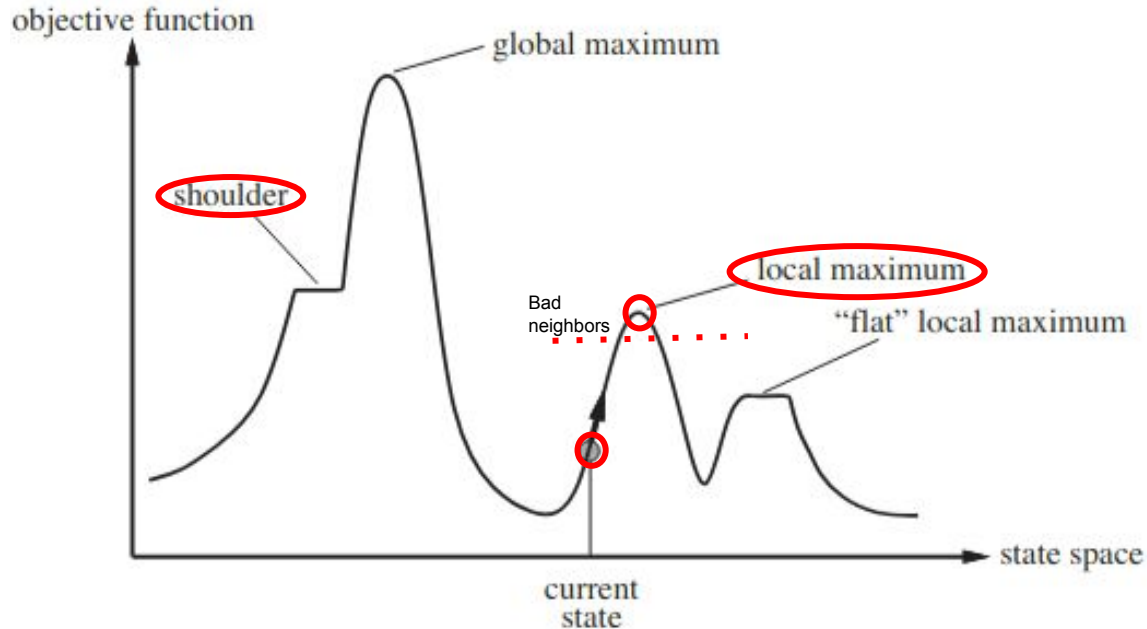
neighbor \leftarrow a highest-valued successor of *current*

if *neighbor*.VALUE \leq *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*

Hill-climbing (Greedy Local Search)


- Is it always possible to reach the global maximum?
 - Depends on where you start.



Hill Climbing gets stuck at a local maxima (or local minima)

Escaping Shoulders: Sideways Move

- If no downhill (or uphill) moves, **allow sideways moves, hoping that the algorithm escapes.**
- Allow movement to **the next state** even if that has **the same objective function value.**
- Need to place a **limit** on the possible **number of sideways moves** to avoid infinite **loops.**
- For 8-queens,
 - Allow sideways moves with a limit of say 100
 - Raises the percentage of problem instances solved from 14% to 94%.
 - However, with this approach, an average of 21 steps are needed for successful solutions and 64 steps for failures.



function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum


```
current ← MAKE-NODE(problem.INITIAL-STATE)
loop do
  neighbor ← a highest-valued successor of current
  if neighbor.VALUE ≤ current.VALUE then return current.STATE
  current ← neighbor
```



```
neighbor ← a highest-valued successor of current
if neighbor.VALUE < current.VALUE then return current.STATE
current ← neighbor
```

Escaping Shoulders: Sideways Move

- Allow movement to the next state, even if that has the same objective function value.
- Problem with allowing Sideways move:
 - May keep oscillating
 - A to B then back to A then B
 - Or oscillate between a set of states



function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

neighbor \leftarrow a highest-valued successor of *current*

if *neighbor*.VALUE \leq *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*




neighbor \leftarrow a highest-valued successor of *current*

if *neighbor*.VALUE < *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*

Escaping Shoulders: Sideways Move

- Allow movement to the next state, even if that has the same objective function value.
- Problem with allowing Sideways move:
 - May keep oscillating
 - A to B then back to A then B
 - Or oscillate between a set of states
- Solution: Reduce the “Amnesia” a little bit and use some memory



function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

neighbor \leftarrow a highest-valued successor of *current*

if *neighbor*.VALUE \leq *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*




neighbor \leftarrow a highest-valued successor of *current*

if *neighbor*.VALUE < *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*


Escaping Shoulders: Tabu Search

- Prevent returning quickly to the same state.
- Maintain a fixed length queue ("tabu list").
- Add the most recent state to a queue.
 - and drop the oldest from the queue in case of full memory.
- **Never make the step that is currently tabu'ed.**
- Properties
 - As the size of the **tabu list grows**, **hill-climbing** will become **like a systematic search** algorithm (basically, don't repeat any state).
 - In practice, a reasonable sized tabu list (around 100) improves the performance of hill climbing in many problems.




Escaping Shoulders: Enforced Hill Climbing

- Perform **BFS** from a local optima
 - To find the **nearest state** with a **better** objective function **value**
 - Effectively adding slightly larger memory
- Typically,
 - Prolonged periods of exhaustive search
 - Bridged by relatively, quick periods of hill-climbing
- Middle ground b/w local and systematic search
- Termination can be tricky -
 - So can limit with time of search



Trivial Algorithms

- Random Sampling
 - Generate a state randomly.
- Random Walk
 - Randomly pick a neighbor of the current state.
- Ignoring memory and time restrictions, these are complete.
- Greedy Local Search/Hill climbing will not necessarily find a solution.
- Can we combine these approaches?



Stochastic Hill-climbing

- Goal: Avoid getting stuck in local minima.
- Strategies:
 - Random selection among uphill moves
 - The selection probability can vary with the steepness of the uphill move.
- Random-walk hill-climbing
- Random-restart hill-climbing
- Hill-climbing with both

Hill-climbing with Random Walk

- Idea: Combine Hill-climbing with Random Walk
- At each step, perform **one of the two**
 - **Greedy** - With probability p move to **neighbour** with the **largest** evaluation function **value** (in case of maximising).
 - **Random** - With probability $1-p$ move to a **random neighbour**.
- Which is more ideal? p constant or variable
 - **Initially**, it may be okay to take **more random** steps. Why?
 - May have started very **close** to **local maxima** (if the goal is maximization).
 - **As time goes on**, we should take more and **more greedy** steps.
 - So **p** should **increase with time**. What happens then?

Hill-climbing with Random Restarts

- Approach:
 - Perform hill-climbing.
 - If you get stuck at a local optima, randomly jump to a new state and perform hill-climbing.
 - Keep the best solution found so far.
- Variations
 - For each restart: run till the greedy optimum and restart, or run for a fixed time and restart.
 - Run a fixed number of restarts or run indefinitely.

Hill-climbing with Both

- At each step, perform one of the three with some probability.
 - Greedy - move to neighbour with the largest evaluation function value.
 - Random Walk - Move to a random neighbour.
 - Random Restart - Resample a new current state.

Simulated Annealing

- Idea
 - **Like hill-climbing**, identify the **quality of the local improvements**.
 - Instead of picking the best move, pick one randomly.
- Say the change in objective function value is δ (= newobjval - oldobjval)
- If δ is positive, then move to that state (when we want to maximize the objective value).
- Otherwise:
 - Move to this state with a probability proportional to δ
 - Therefore, worse moves (very large negative δ) are executed less often.
- However, there is always a chance of escaping local maxima.
 - Therefore, overtime, make it less likely to accept locally bad moves.

Simulated Annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

inputs: *problem*, a problem

schedule, a mapping from time to “temperature”

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

for $t = 1$ **to** ∞ **do**

$T \leftarrow \text{schedule}(t)$

if $T = 0$ **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow \text{next.VALUE} - \text{current.VALUE}$

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$

The algorithm shows that the simulated annealing algorithm is a **version of stochastic hill climbing** where **some downhill moves are allowed (in case the objective is to maximize the objective function value)**.

- **Downhill moves** are **accepted** readily **early** in the annealing schedule and then **less often** as time goes on.
- The **schedule input** determines the **value of the temperature T** as a function of time.
- Symbolic meaning of T:
 - If the **temperature** is **low**, we are stable and can decide properly based only on **greedy** steps.
 - If the **temperature** is **too low**, we are frozen and can't move, then the **current state is the solution**.
 - If the **temperature** is **high**, we are unstable and may make bad decisions, so **allow downhill moves**.

Simulated Annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

inputs: *problem*, a problem

schedule, a mapping from time to “temperature”

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

for $t = 1$ **to** ∞ **do**

$T \leftarrow \text{schedule}(t)$

if $T = 0$ **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow \text{next.VALUE} - \text{current.VALUE}$

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$

Temperature T

- high T: probability of “locally bad” move is higher.
- low T: probability of “locally bad” move is lower.
- Typically, T is decreased as the algorithm runs longer.
 - i.e., there is a temperature schedule



Physical Interpretation of Simulated Annealing

- A Physical Analogy:
 - imagine letting a ball roll downhill on the function surface
 - this is like hill-climbing (for minimization)
 - now imagine shaking the surface, while the ball rolls, gradually reducing the amount of shaking
 - this is like simulated annealing

Local beam search

- Idea: Keeping only one node in memory is an extreme reaction to memory problems
- Keep track of k states instead of one [Beam of size k]
 - Initially: k randomly selected states
 - Next: determine all successors of k states
 - If any of the successors is a goal then finish
 - Else select k best from successors and repeat

Local beam search

- Not the same as k random-start searches run in parallel!
- Searches that find good states recruit other searches to join them
- Problem: quite often, all k states end up on same local hill
- Idea: Stochastic beam search
 - Choose k successors randomly, biased towards good ones
 - better valued states have higher probability of being selected
 - lesser valued states may also get selected but with a lower probability
 - Improves local beam search by introducing a degree of randomness into the selection process, maintaining diversity among candidate solutions, and reducing the risk of premature convergence to local optima.
- Observe the close analogy to natural selection!